

Brandon Willett

Education

UNIVERSITY OF ROCHESTER

- ◇ Graduated in 2018
- ◇ B.S. in Computer Science (3.82)
- ◇ B.A. in Mathematics (3.80)

Skills

BACKEND

- ◇ Go, Python, Java, Scala, Javascript
- ◇ Web Frameworks (Django, Flask, Rails)
- ◇ Streaming Pipelines (Kafka)
- ◇ Workflow Platforms (Airflow, Luigi)
- ◇ Transactional Databases (Postgres, Cass)
- ◇ Analytical Databases
- ◇ Caches (Memcached, Redis)

PLATFORM

- ◇ Version Control (Git)
- ◇ CI/CD (Github Actions, Gitlab)
- ◇ Telemetry (Prometheus, Datadog)
- ◇ Containerization (Docker, Containerd)
- ◇ Build Tools (Rake, Bazel)
- ◇ Infra as Code (CF, Terraform)
- ◇ Service Discovery (DNS, Consul, Envoy)
- ◇ Load Balancers (L4, L7)
- ◇ Reverse Proxies (NGINX, HAProxy)

CLOUD

- ◇ AWS (ALB, NLB, EC2, RDS, S3, Lambda)
- ◇ GCP (GLB, GCE, GCS, Cloud SQL)
- ◇ Azure (ALB, VMSS, ABS)
- ◇ Kubernetes (ask me about any resource)

I'm Brandon, nice to meet you! I'm a versatile developer, experienced in large-scale data processing, API serving, and cloud infrastructuring. I equally enjoy working with good software and good people, and I like helping build collaborative cultures.

☎ +1 203 258 2721

✉ brandon@willett.io

Experience

TODOIST

Remote

Backend Engineer

Mar 2024 – Present

- ◇ Worked on improvements to our public-facing APIs and sharded MySQL
- ◇ Noticed and preemptively fixed several architectural issues over what was supposed to be just a run-of-the-mill Elasticsearch migration
- ◇ Held sessions around observability best practices for dev and CX teams

DATADOG

New York, NY

SRE / Platform Engineer

May 2020 – Jun 2023

- ◇ Worked in network edge and metrics platform to build new systems, from design to migration to on-call, and maintain existing ones
- ◇ Identified and prioritized consistently high-ROI work, like a system which delivered configuration to 10+ services simultaneously via our pre-existing Consul (allowing breathing room to write a real RPC replacement)
- ◇ Designed and got consensus on a Kafka architecture, with smart partition choices, allowing metrics to scale beyond its primary "god-cluster"
- ◇ Completed adoption of the new Kubernetes platform in metrics, migrating some of Datadog's largest + most complicated services (the primary TSDB, the primary public LB) without downtime
- ◇ Worked across boundaries, fixing the largest bottlenecks (in both software and people-process) regardless of where the change needed to be
- ◇ Hosted architecture reviews and mentored lots of new hires, to help folks feel empowered, instead of confused, as quickly as possible

ACTIONIQ

New York, NY

Backend Engineer

Jun 2018 – Apr 2020

- ◇ Designed, advocated for, and implemented the procedure for autoscaling query execution on EC2, reducing our AWS bill by more than 50%
- ◇ Drove scalability improvements for the OSS job scheduling platform Luigi, and contributed some of that back upstream
- ◇ Maintained the Prometheus cluster (later, Datadog) and led internal talks and workshops on best practices around metrics and observability
- ◇ Brought some of the first Terraform modules to AIQ, used them to build the first version of our streaming ingest platform using IaC