

Brandon Willett

Education

UNIVERSITY OF ROCHESTER

- ◊ Graduated in 2018
- ◊ B.S. in Computer Science (3.82)
- ◊ B.A. in Mathematics (3.80)

Skills

BACKEND

- ◊ Go, Python, TypeScript, Scala
- ◊ Web Frameworks (Django, Flask, Rails)
- ◊ Streaming Pipelines (Kafka)
- ◊ Workflow Platforms (Airflow, Temporal)
- ◊ Transactional Databases (Postgres)
- ◊ Analytical Databases
- ◊ Caches (Memcached, Redis)

PLATFORM

- ◊ Version Control (Git)
- ◊ CI/CD (GitHub Actions, Gitlab)
- ◊ Telemetry (Prometheus, Datadog)
- ◊ Containerization (Docker, Containerd)
- ◊ Build Tools (Rake, Bazel)
- ◊ Infra as Code (CDK, Terraform)
- ◊ Service Discovery (DNS, Consul, Envoy)
- ◊ Load Balancers (L4, L7)
- ◊ Reverse Proxies (NGINX, HAProxy)

CLOUD

- ◊ AWS (ALB, NLB, EC2, RDS, S3, Lambda)
- ◊ GCP (GLB, GCE, GCS, Cloud SQL)
- ◊ Azure (ALB, VMSS, ABS)
- ◊ Kubernetes (ask me about any resource)

I'm Brandon (👋), a versatile backend and infra developer who enjoys working on honest software with honest people.

☎ +1 203-258-2721 ✉ brandon@willett.io

Experience

CURSORM (via GRAPHITE)

SRE / Infra Engineer

New York, NY

Apr 2025 – Jul 2026

- ◊ Architected a complete DB reliability platform with query cost-limiting and connection pooling, drastically reducing DB outages and timeouts
- ◊ Owned the entire on-call system, not only creating dozens of monitors but also implementing seamless alert routing for split rotations
- ◊ Designed a clever code index with Turbopuffer for custom MCP tools
- ◊ Orchestrated several complex migrations with no (or minimal) downtime, like domain (.dev → .com) and AWS region (us-west-2 → us-east-1)

TODOIST

Backend Engineer

Remote

Mar 2024 – Mar 2025

- ◊ Identified and fixed several critical architectural issues preemptively during a major Elasticsearch migration
- ◊ Held sessions for observability best practices for dev and CX teams

DATADOG

SRE / Platform Engineer

New York, NY

May 2020 – Jun 2023

- ◊ Drove full-lifecycle system development in the network edge and metrics platforms, crossing organizational boundaries to eliminate bottlenecks in both software and people-process
- ◊ Completed the zero-downtime migration of Datadog's largest and most complex services (primary TSDB, primary public LB) to Kubernetes
- ◊ Kicked off several reliability initiatives, including founding a group to lead weekly metrics-wide incident learnings, unified SLI tracking, and running personally many production-readiness reviews
- ◊ Designed and got consensus on a Kafka architecture, with smart partition choices, allowing metrics to scale beyond its primary "god-cluster"

ACTIONIQ

Backend Engineer

New York, NY

Jun 2018 – Apr 2020

- ◊ Designed, advocated for, and implemented the procedure for autoscaling query execution on EC2, reducing our AWS bill by more than 50%
- ◊ Created and contributed scalability improvements to Luigi (open-source)
- ◊ Maintained the Prometheus cluster (later, Datadog) and led internal talks and workshops on best practices around metrics and observability